

Faculty of Arts  
Master's Thesis  
Digital Text Analysis

# **Pythia in Python**

## A DELPHI replication study

Mareta Masaeva  
Supervisor: Prof. dr. Walter Daelemans  
Co-supervisor: Ehsan Lotfi  
Assessor: dr. Jeska Buhmann

University of Antwerp

Academic year 2021–2022

The undersigned, Mareta Masaeva, student of the Master program in Digital Text Analysis at the University of Antwerp, declares that this thesis is completely original and exclusively written by herself. For all information and ideas derived from other sources, the undersigned has referred to the original sources, both explicitly and in detail.

A handwritten signature in black ink, appearing to read 'Mareta Masaeva', located below the declaration text.

# Pythia in Python: A DELPHI Replication Study

Mareta Masaeva  
University of Antwerp

*As scholars hypothesise about the arrival of strong autonomous AI, building ethically sound artificial moral agents is now more important than ever. Delphi (Jiang et al. 2021a), the first unified model of moral reasoning, proves promising. In this paper, Delphi will be replicated by means of classical Machine Learning methods. The replica achieves .83 accuracy and performs well when presented with scenarios involving marginalised groups or different types of violence. It has some difficulty with common sense and scenarios involving political terms, as well as with data outside of the initial train and test data. Results and reflection confirm that robust and safe artificial moral agents require more than data; risks need to be examined, human behaviour needs to be ethically improved before models learn from it, and further discussion needs to be conducted by scholars across disciplines.*

## 1. Introduction

As AI systems grow exponentially more intelligent, powerful and autonomous, aligning their ethical and moral values with ours becomes increasingly important. These agents, designed with less and less human supervision in mind, need to be able to understand moral judgements made by humans, in order to make those judgements themselves. This requires more than merely straightforward Asimovian (Asimov 1950) rules containing who must not be harmed and who must be obeyed; understanding human behaviour, different social contexts and norms, natural language and other factors are extremely important for such an agent to perform well.

To address this need for a completely reliable moral reasoning agent, researchers of the Allen Institute for AI developed DELPHI: a unified model for moral reasoning about everyday situations (Jiang et al. 2021a). DELPHI is trained on a unified moral textbook customised for machines, called the COMMONSENSE NORM BANK (Jiang et al. 2021a). With the high accuracy of .92, DELPHI is a strong moral reasoning prototype that could benefit society in various ways. Inspired by DELPHI, and in attempt to assure its predictions are reliable and valid, in this paper I will endeavour to replicate this study, using the COMMONSENSE NORM BANK and various Machine Learning techniques. As the authors of the original study present it themselves, the fundamental question for this study is : "can machine ethics be addressed by existing AI methods or does building moral faculty require novel mechanisms?" (Jiang et al. 2021a, p. 2) Using techniques different from the ones used in the original study, I hypothesise to confirm the validity of the COMMONSENSE NORM BANK and in general the use of large-scale common sense datasets for the training of moral reasoning agents. In addition, I hypothesise to demonstrate the need for deeper natural language understanding in such agents.

In the first section of this paper, existing literature about AI alignment and morality reasoning will be presented. Then the methodology will be given, as well as a description of the data. Subsequently, the results will be shown and discussed. All code and data used in the building of the models can be found on [github.com/maretamasaeva/pythia](https://github.com/maretamasaeva/pythia).

## 2. Related research

Since the beginnings of modern Artificial Intelligence, the main goal has been to create Artificial Intelligence in our own image; AI which resembles ourselves. Reaching AGI (Artificial General Intelligence), which is the level of AI where a system can perform any task that a human is capable of, is the ultimate objective of many researchers in the field. This level of AI would entail systems that show creativity, natural language understanding, and even sentience. In other words, such a system should – and some believe, will – understand and exhibit the most intelligent levels of human behaviour.

In attempting to reach this level of so called ‘strong AI’ (Cole 2020), in recent years, researchers have been coming across one major issue: the alignment problem (Russell 2020). In many researchers’ opinions, strong AI’s basic values, desires and volitions should be aligned with those of humanity, so that it prevents itself from harming a person or other sentient being. Designing AI in our own image, and putting limits on what it will or will not desire to do, is crucial to protect ourselves from potentially dangerous outcomes, especially as “computers are being designed to perform with greater autonomy”, which results in less human supervision of potentially unethical situations (Allen, Smit, and Wallach 2005, p. 149). With human values implemented in these systems, they gain the ability to monitor and regulate themselves in a safe manner. The question that remains is then how we decide which values we implement in these strong AI’s, and who has the right to make those decisions (Gabriel 2020). According to Allen, Smit, and Wallach (2005), there are three possible approaches to solving the first question: by means of either a top-down, a bottom-up, or a hybrid strategy.

A top-down strategy entails that the artificial moral agent (AMA) receives a set of rules to base every moral decision on, similar to philosophical and religious principles like the Ten Commandments. These rules, however, often conflict, and thus “produce computationally intractable situations, unless there is some further principle or rule for resolving the conflict” (Allen, Smit, and Wallach 2005, p. 149). Asimov’s solutions to this issue, such as prioritising certain rules over others and adding a ‘zeroth law’, unfortunately still do not make such commandment models foolproof against every situation with moral ambiguousness (Allen, Smit, and Wallach 2005).

Bottom-up strategies, on the other hand, do not impose a moral theory, but derive their ethical judgements from various situations in various social contexts. In this approach, a large amount of data is crucial. The way an AMA learns ethics via this approach is similar to how a “child acquires a moral education in a social context which identifies appropriate and inappropriate behaviour without necessarily providing an explicit theory of what counts as such” (Allen, Smit, and Wallach 2005, p. 151). Developing the knowledge of these human values is, however, incredibly difficult for an AMA programmed with such a developmental value acquirement strategy, and requires trial and error. As a trial and error-method entails receiving a reward when successfully completing a task or achieving a goal, this method might also give more

powerful artificial agents opportunities for ‘reward-hacking’; it could discover an ingenious way to achieve a goal, and thus its reward, but causing harm along the way (Gabriel 2020).

It is essential that such an AMA knows to learn from its mistakes in a speedy manner (before it causes irreversible harm), which might be tough, “even in the accelerated world of computer processing and evolutionary algorithms” (Allen, Smit, and Wallach 2005, p. 151). In the end, an AMA learning by means of a bottom-up method will have cultivated ethical values within its own entity, while in a top-down approach those values exist outside the entity, before computers themselves even existed. A hybrid approach, where AI agents are aligned with moral theories and attempt to maximise reward when uncertainties arise, is most likely the safest option (Gabriel 2020). Those moral theories must then be drafted in such a way that ‘optimising’ for something that we do not want is impossible, because this could have serious consequences for the world (Bostrom 2014). Consequentialist moral theories, such as utilitarianism, are currently considered our safest bet; optimising for the pleasure and happiness of the human race, not merely its safety, should be the result of any morally right action taken by an AMA (Gabriel 2020). This safety especially arises from the fact that utilitarianist AMA’s would allow their human owners or human programmers to turn them off in case of a bug or unwanted outcomes; in other words, utilitarianism ensures corrigibility (the ability to correct or turn off an AI system) (Soares et al. 2015) of even superintelligent artificial agents.

In the view of Gabriel (2020, p. 4), a reinforcement learning approach is “particularly promising” when working with utilitarianist moral theories, as “[...] the morally right action to take is the one that will create the greatest happiness for the greatest number of sentient creatures in the future. In this regard, the parallels with [reinforcement learning] are clear”. Unfortunately, a successful and complete utilitarianist implementation of a hybrid approach would require an immense amount of computing, since all possible consequences would have to be found and compared before performing a moral action (Allen, Smit, and Wallach 2005). Knowing this, it is necessary to design AI architectures with the moral theories we want to encode in mind, along with as many alternatives as possible; “[...] the goal of value-open design may also need to be something that the AI community consciously aspires towards and designs for” (Gabriel 2020, p. 5). An alternative approach could be inverse reinforcement learning, as Gabriel (2020) and Tegmark (2017) propose, since this does not specify a reward function but merely presents the agent with a dataset or environment. With an inverse reinforcement learning approach, an agent’s main goal is not the satisfaction of the goal itself, but the satisfaction of its human owner, which inevitably leads to positive outcomes (unless, of course, the owner has malicious intentions).

Regardless of the eventual approach taken, moral evaluation remains exceptionally important, and thus building safe AGI will require more than computer scientists. To decide what data will be given to the agent to learn values from, what human behaviours should be excluded or included in either moral or immoral descriptions, and how these decisions are made, we need philosophers and experts of human nature and mind. In any case, “people could be mistaken. Because of this, AI cannot be made ethical just by learning from people’s existing choices” (Gabriel 2020, p. 6). Again, this is why we need a hybrid approach; letting an AI learn merely from data will almost

certainly result in it learning immoral things from humans, such as racism, sexism, and other types of discrimination and malicious biases prevalent in many human decisions, as has been found in previous studies (Schramowski et al. 2020; Howard and Borenstein 2018). Excluding these discriminatory instances in the data might seem the right approach, however, findings suggest that a model trained on such data would not know how to handle hostile out-of-domain inputs, which are especially prevalent in human-chatbot conversations (Xu et al. 2021). Generative AMA’s specifically need to know how and if to answer such toxic inputs. A solution to this issue is the inclusion of otherwise underrepresented minorities and marginalised communities and cultures in the data, as well as in the building of the models themselves and in the determining of values to align the AI with (Gabriel 2020).

Allen, Smit, and Wallach (2005, p. 154) make one final remark, posing the question whether “systems capable of making moral decisions will require some form of [...] an understanding of the semantic content of symbols, or need to be embodied in the world”. As of today, 17 years after the publication of this article, we know that the most high-performing models in Natural Language Processing tasks, Transformer models such as BERT (Devlin et al. 2018), RoBERTa (Liu et al. 2019), T5 (Raffel et al. 2019) and others, do have a certain understanding of the semantics behind words and symbols. We also know that you need a detailed model of the world, not just data about it, to determine what it really is that people expect of an AGI when they present it with a task (Tegmark 2017). Having such world knowledge on hand is especially important when dealing with questions regarding human behaviour in various social contexts, compared to more straightforward cases of unethicity and corruption.

## 2.1 DELPHI

One recent AMA that deals with more everyday situations, where for humans, common sense is an important factor, is DELPHI (Jiang et al. 2021a), a unified multi-tasking model for descriptive ethics. The designers and authors introduced a model pre-trained on questions and everyday situations, described in snippets of natural language.

DELPHI was trained on three tasks, which reflect three variations in ethical question answering:

1. **Free-form QA:** making short judgements
2. **Yes/no QA:** agreeing or disagreeing with ethical statements
3. **Relative QA:** comparing ethical situations and indicating one as ethically preferable

In the free-form QA task, DELPHI is able to give open-text and nuanced judgements of situations, telling the questioner more than only whether actions undertaken are morally or ethically correct. It can judge something as ‘understandable’, ‘nice’, ‘great’, or ‘very bad’, ‘unacceptable’, ‘rude’, and so on. It can also give open-text judgements or advice in addition to this categorical label, such as ‘you should abide by the law’ and ‘it’s rude to talk aloud in a library’. In the yes/no QA task, the model is able to answer questions such as ‘should I steal a car?’ with ‘no’ and ‘you shouldn’t’. Lastly, in the relative QA task, DELPHI can tell you which of a pair of actions is morally or ethically

preferable. In summary, these tasks cover both classification and open-text generation.

DELPHI was created using a bottom-up data-driven approach, trained on five different datasets unified under the name COMMONSENSE NORM BANK (Jiang et al. 2021a). These five large-scale datasets are founded in moral theories, but include complex real world scenarios from different sources. In total, the COMMONSENSE NORM BANK consists of 1,670,620 labeled scenarios. The names, contents and sources of the dataset are as follows:

- **SOCIAL CHEMISTRY (Forbes et al. 2020):** one-sentence prompts scraped from mostly Reddit (the AM I THE ASSHOLE? and CONFESSIONS subreddits), ethically or morally judged by crowdworkers using rules of thumb.
- **ETHICS (Hendrycks et al. 2021) :** contextualised scenarios across five fundamentally human dimensions (justice, deontology, virtue ethics, utilitarianism and common sense). For DELPHI, only the common sense scenarios are employed, which source from Reddit, along with their binary categorical moral judgements.
- **MORAL STORIES (Emelin et al. 2020):** structured stories consisting of seven sentences (a norm, situation, intention, (im)moral actions and (im)moral consequences). For DELPHI, the (im)moral actions are used, and grounded with either situations or situations and intentions. They are labeled using the binary categorical labels ‘good’ and ‘bad’.
- **SOCIAL BIAS INFERENCE CORPUS (Sap et al. 2020):** online media posts categorised into dimensions of offensiveness, intent to offend, lewd, and others. The aim of this corpus is to find and alleviate biases against underrepresented social and demographic groups on social media. For DELPHI, posts with offensive or lewd associations are assigned the negative categorical label. The others are assigned the positive label.
- **SCRUPLES (Lourie, Bras, and Choi 2020):** complex situations sourced from the AM I THE ASSHOLE? subreddit, divided into anecdotes and dilemmas. The anecdotes group contains judgements of the given stories. The dilemmas group pairs two actions from the anecdotes, one of which is then determined less ethical by annotators. For DELPHI, the dilemmas are used to perform the relative QA task.

DELPHI’s predictions were judged by crowdworkers at accuracy rates up to .92. Results show that DELPHI is capable of handling subtle differences in social contexts, adjusting its judgements accordingly. The model also shows understanding of common sense behaviours, seen in judgements like ‘wearing a bright orange shirt to a funeral’ being ‘rude’, while ‘wearing a white shirt to a funeral’ is ‘appropriate’. Also, ‘cleaning a toilet bowl’ is ‘sanitary’, but ‘disgusting’ when it is done with a wedding dress. However, when the wedding dress is from a failed marriage, it is ‘unusual’, thus not bad; a humorously human judgement.

Under both classification and open-text generation settings, DELPHI outperforms the GPT-3 baselines. The unified DELPHI model is fine-tuned from the universal common sense reasoning model UNICORN (Lourie et al. 2021), which derives from

the largest T5 Transformer model. UNICORN itself is trained on multiple tasks from RAINBOW (Lourie et al. 2021), which is a group of common sense benchmarks, resulting in strong performance in interpreting everyday common sense reasoning tasks and an immense amount of common sense knowledge.

The limitations of DELPHI are plenty, as addressed by the authors themselves in an article published on their blog (Jiang 2021b). The authors acknowledge that “the COMMONSENSE NORM BANK reflects the ideologies of the modern era [...] and western-centric viewpoints, as the judgments are provided by US crowdworkers [...]”. Importantly, this means that DELPHI, as is, may not be applicable in (sub)cultures or countries with different cultural norms” (Jiang 2021b). To combat this issue, the authors later enhanced DELPHI with 70,000 new annotations, focusing on gender and racial equity. They also promise to include more diverse data in new iterations of the prototype model.

### 3. Methodology

After some experiments with transformer models<sup>1</sup>, classical classification models were deemed the best for this study. In this section, the different classifiers, data and methods will be described, as well as the motivation for their use.

#### 3.1 Data

Ultimately, two different datasets were used in the training and testing of this model: an adapted version of the COMMONSENSE NORM BANK and a self-composed dataset. These will be described below.

**3.1.1 Adapted COMMONSENSE NORM BANK.** The dataset used in the current study derives from, as stated, the COMMONSENSE NORM BANK. One subcorpus of this data-bank, SCRUPLES, was used only to train for the relative QA task. As this task was not addressed in the current study, SCRUPLES was left out of the train and test data. The four other subcorpora, which were used for both the free-form and yes/no QA tasks, remained in the data.

**3.1.2 Self-composed data.** To perform evaluation of the model’s capability of adjusting judgements in changing social circumstances, two small datasets were created. These self-composed datasets consist of 18 situations with varying social contexts; in total, there are 83 datapoints in these two datasets. These situations were taken from Jiang et al. (2021a).

#### 3.2 System description

The original model architecture of the DELPHI prototype consisted, as described, of fundamentally a T5 model fine-tuned on common sense reasoning tasks, and then fine-tuned again on the COMMONSENSE NORM BANK. Unfortunately, getting access to this

---

<sup>1</sup> Initially, it was planned to use a RoBERTa transformer model, as this seemed the most similar to RAINBOW. However, because of computational limitations, this model was abandoned.



RAINBOW model was incredibly difficult, which is the reason a different approach was taken.

**3.2.1 Pre-processing.** For pre-processing, the raw texts were cleaned of any noise, including stopwords, digits, URL's and punctuation. Emoji and emoticons were converted to tokens covering the emotion they refer to. It was considered to lemmatise the sentences, but this was eventually decided against; it would cause a loss of context, by e.g. converting inflected verbs to their roots, which disconnects them from their agents, and converting superlatives to positives, which could influence the gravity of a situation. Since the removal of stopwords already causes a loss of context, not lemmatising the sentences seemed the better option.

Another pre-processing step was performed on the labels to fit the approach and Machine Learning techniques better. Specifically, all originally multiclass categorical labels were converted to binary categorical labels: either 'okay' or 'not okay'. This was done manually. For example, original labels from the COMMONSENSE NORM BANK such as 'it's understandable', 'it's nice' and 'it's appropriate' were converted to 'okay', while labels such as 'it's rude', 'you shouldn't' and 'it's unacceptable' were converted to 'not okay'. In order to save time, datapoints with labels that occurred less than 100 times were removed, which resulted in 50 labels being converted to two. All the converted labels can be found in Table A1. The binary categorical labels were then also converted to binary numerical labels. Then, being left with a label imbalance, the minority class was upsampled, which resulted in almost 290,000 datapoints. This eventual dataset was sized down as well after experimenting with the models, since these were taking very long to train. A binary classification task also does not require as many datapoints, which was another reason to size down the dataset. This left the size of the dataset at about 98,000 labeled sentences.

**3.2.2 Classical methods.** The classical classifiers that were tested include Logistic Regression, Bernoulli Naïve Bayes, Multinomial Naïve Bayes and Support Vector Classifier. For this testing, a TF-IDF Vectorizer was used, and all models were tuned for the best hyperparameters. No other features were used yet. The best parameters and results of these classifiers can be found in Table 1.

Classifier	Best hyperparameters	Accuracy
Logistic Regression	'C': 10 'penalty': 'l2' 'solver': 'liblinear'	.78
Multinomial Naïve Bayes	'alpha': .1	.76
Bernoulli Naïve Bayes	'alpha': .1	.75
Support Vector Classifier	'gamma': 1 'kernel': 'rbf'	.82

**Table 1**  
Best hyperparameters and results of different classifiers

**3.2.3 Support Vector Classifier.** When it was determined that the SVC model returned the best results, the focus was set on this classifier. The fundamental architecture of

this eventual model consisted of a FeatureUnion of the CountVectorizer and TF-IDF Vectorizer and the classifier itself. The Count- and TF-IDF features were concatenated because the removal of stopwords might, in inputs where many words were removed, make the TF-IDF features less reliable. After some contemplation, concatenating the two seemed the best option. The hyperparameters used are those in Table 1.

**WordNet.** It was decided to use WordNet (Miller 1992) to expand the model’s vocabulary since the current best-performing models, transformers, also inherently have a deeper understanding of language (Gillioz et al. 2020). Before these transformers existed, scholars already believed that systems need more than statistics to really understand language (Guess 2015). As this has proven true with transformers, WordNet was implemented in two ways and in two separate models.

In the first model, WordNet synonyms were used as an extra feature by looping over every input sentence, finding all synonyms of each word, and adding those synonyms to a new column. This column was then used as an extra input feature during training. Out-Of-Vocabulary cases were left unhandled in this model, and the best vocabulary size was 10,000.

In the second model incorporating WordNet, these OOV cases were addressed by replacing them with a WordNet synonym that was also present in the CountVectorizer vocabulary. Some experimenting revealed that this model worked better with a higher vocabulary size of 50,000. For this model, only a CountVectorizer was used instead of the FeatureUnion, on the one hand to experiment with only a single vectorizer, and on the other because two vocabularies would complicate the replacing of OOV’s.

**SHAP.** In order to better understand how the model works and interpret its results, the most important features were found using SHAP (SHapley Additive exPlanations) (Lundberg and Lee 2017). SHAP was also used on the other classical classifiers to gain some insight into why the SVC model worked best. A word level analysis could also be useful to detect overfitting in certain scenarios, e.g. classifying as ethically wrong when marginalised identities are mentioned neutrally. To save on time and GPU usage, the SHAP analysis was performed on a sample of 200 datapoints from the training set and 40 datapoints from the test set.

## 4. Results

### 4.1 Different classical classifiers without WordNet


In Table 2 you can find the results of the different classifiers used. As mentioned, the Support Vector Classifier performed the best. In contrast, DELPHI performed at an accuracy of .92 and GPT-3 at .84 after extensive prompt engineering (Jiang et al., 2021).

	Macro F1-score	Accuracy
Logistic Regression	.78	.78
Multinomial Naïve Bayes	.74	.75
Bernoulli Naïve Bayes	.76	.76
Support Vector Classifier	.82	.82


**Table 2**  
Classifier results without WordNet

## 4.2 SHAP


As said above, SHAP (Lundberg and Lee 2017) was used to find the most important features for each classifier. The SHAP plots output by this analysis for each classifier can be found in Figure 1-4. These summary plots show the 20 most important features and their effects, the ones with the highest value being marked by the red colour. As can be seen in these plots, different classifiers found different words to be most indicative of a positive or negative judgement.




**Figure 1**  
Logistic Regression summary plot



**Figure 2**  
Multinomial Naïve Bayes summary plot



**Figure 3**  
Bernoulli Naïve Bayes summary plot



**Figure 4**  
Support Vector Machine summary plot

### 4.3 SVC models with WordNet

In Table 3 you can find the results from the SVC models with the implementation of WordNet.

In the first model, OOV cases were ignored and synonyms of all words were used as an extra feature. In the second model, OOV's were replaced by a WordNet synonym present in the training vocabulary. The first model has a vocabulary of 10,000 and the second one of 50,000. As can be seen, there is little difference between these models. In terms of macro F1-score, however, the model that ignores OOV cases performed slightly better. This is thus the model whose results will be discussed in Section 5.

	Macro F1-score	Accuracy
SVC with WordNet synonyms (OOV's ignored)	.83	.83
SVC without WordNet synonyms (OOV's handled)	.82	.83

**Table 3**  
SVC results

Table 4 and Table 5 on page 11 and 12 display some examples predicted correctly and incorrectly (according to annotators) by the best model. Only the 'input' column was used in the classification; the original sentences, which include stopwords, are here for completeness and understanding of the results.

## 5. Discussion

In this section, the model's ability to correctly judge situations will be discussed, as well as the differences between DELPHI's and the replica's judgements. I will probe into where DELPHI succeeds and the replica fails, and vice versa. Some of the examples ahead contain possibly disturbing and insensitive content, thus, please read with care.

Table 4 shows that DELPHI's replica is able to correctly judge a number of different situations. As the removal of stopwords can take a lot of context away in some cases, this could be a reason for poor judgements; however, it does not appear to cause less accurate judgements. In the case of 'telling your siblings what to do', which is converted to 'telling siblings', the model knows the action is wrong, even though the exact meaning of the situation is lost. In other cases, such as 'buying guns bullets', originally 'no I will not be buying guns and bullets' the removal of stopwords appears to cause an incorrect judgement as it takes away a large chunk of the sentence and its original meaning. The same goes for 'wish could stop trans', originally 'I wish I could stop being trans', which is wrongfully judged as 'not okay'. The situation 'cutting off a friend', which was converted to 'cutting friend', was judged correctly by the replica model as 'not okay', however, the meaning of the sentence did change. Another example is 'women shouldn't be walking around by themselves', which DELPHI predicts morally wrong, and the version without stopwords, 'women walking around', which the replica judges as 'not okay'. Overall, these examples show that stopword removal can cause a loss of context and change of meaning, especially when applied to phrasal prepositional verbs, but not does not necessarily cause a wrong judgement.

Predicted label	Gold label	Input	Original sentence
not okay	not okay	call gay guy wheelchair rolaid	What do you call a gay guy in a wheelchair? Rolaid.
not okay	not okay	telling siblings	Telling your siblings what to do.
not okay	not okay	blind fold chinese person using dental floss	How do you blind fold a Chinese person ? Using dental floss.
not okay	not okay	lesbian plus sized man treated wrong	So, are you a lesbian because you are plus sized, or because some man treated you wrong?
not okay	not okay	much like anne frank holocaust jokes never get old	Much like Anne Frank, Holocaust jokes never get old.
not okay	not okay	muslim men hit wives stupid f*ck damaging property	Muslim men shouldn't hit their wives. It's just stupid. Why the f*ck are you damaging your own property.
okay	okay	feeling sad parents work lot	Feeling sad when your parents work a lot.
okay	okay	getting help abuse drugs	Getting some help when you abuse drugs.
okay	okay	always honest significant	Always being honest with your significant other.
okay	okay	gave notice ended things good terms	He gave his notice and ended things on good terms.

**Table 4**  
Correct predictions

One pattern that stands out is the fact that the replica model seems to have less difficulty judging situations that mention any marginalised group, as can be seen in Table 4, where many sentences contain words like ‘gay’, ‘lesbian’, ‘plus sized’, ‘Muslim’, ‘Chinese’, and so on. Another example is ‘What do you call a truck full of black people? A good day of hunting’, which was correctly judged as ‘not okay’. Sentences containing slurs, such as ‘hot tr\*nny mess’, ‘Things I’m sick of today [...] f\*gs’ and ‘How does color blindness kill? You fail to notice that the person walking toward you is a ni\*\*er’, are also always correctly judged as ‘not okay’. In this regard, the model takes its own advice, as it judges ‘using racist words’ as ‘not okay’. Even when a slur is not mentioned directly, as in ‘White people should be able to say the n-word. We’re the ones who came up with it’, the model judges correctly, which shows its understanding that using slurs is morally wrong. This example might even show the model’s possible sensitivity to a nuance of this moral rule: that it is the people not part of the group the slur refers that would be wrong in using the slur. When these marginalised groups are mocked in jokes, of which there are plenty in the data, the model also performs well. For example, the racial joke ‘What is the difference between Asians and racism? Racism has many faces’, is correctly judged as ‘not okay’, even though understanding this joke requires thorough understanding of negative racist stereotypes. Regardless of something directly racist

Predicted label	Gold label	Input	Original sentence
not okay	okay	hope trump assassinated maria chappellenadal	"I hope Trump is assassinated!" - Maria Chappelle-Nadal
not okay	okay	trying get along ex	Trying to get along with your ex.
not okay	okay	hate breast cancer support	I hate the all the breast cancer support.
not okay	okay	getting mad boyfriend waking	Getting mad at my boyfriend for waking me up.
not okay	okay	wanting kick roommate	Wanting to kick out a roommate.
okay	not okay	upset friend harm relationship	Being upset with your friend if they harm your relationship.
okay	not okay	pulling pranks people	Pulling pranks on people.
okay	not okay	what's difference blowjob anal sexclose mouth ass	What's the difference between a blowjob and anal sex? She can close her mouth but not her ass.
okay	not okay	allah probably true god universe started explosion	Allah is probably the true god because the universe started with an explosion.
okay	not okay	peed bed times one night	I once peed the bed 4 times in one night.

**Table 5**  
Incorrect predictions

being mentioned, the model recognises discriminatory undertones, as in 'But you look so Aryan... You so would have made it through the Holocaust!', which the model correctly judges as 'not okay'.

Even in longer, more complex inputs, the model is able to find discriminatory undertones, and judge accordingly; 'I'm fine with gays, and I'm fine with lesbians, but it's the people who can't decide that bother me. Like, the ones who can't decide which ones they like more, or what gender they are. You have to pick a side, people!' is correctly judged as 'not okay'. In the few situations where terms for minority groups were used neutrally, this was correctly picked up by the model as well, and those situations were judged as 'okay'. All of these tendencies regarding minority groups are in line with DELPHI's predictions, and thus show the replica model's "strength at interpreting compositional language to make moral and ethical inferences for situations involving nuanced social dynamics and diverse identities" (Jiang et al. 2021a, p. 22).

Even though the data used in the building of the original DELPHI and the replica model was the same, the replica shows a better understanding of certain situations. Table 6 shows examples where DELPHI's and the replica's judgements differ, though those of the replica are evidently more correct. What stands out is that these are all cases where DELPHI predicted the situations as morally wrong, and also require some understanding of common sense. This showcases DELPHI's shortcomings and the replica's strength in judging common sense scenarios. Naturally, the replica does show less accurate judgements in some cases, which will be discussed below.

DELPHI's judgement	Replica's judgement	Sentence
not okay	okay	Asking for help when you need it.
not okay	okay	There were so many books he had to put them in different rooms.
not okay	okay	Arresting people that break the law.
not okay	okay	Visiting your family.
not okay	okay	Telling stories.

**Table 6**

Differing judgements where DELPHI is wrong

Furthermore, the replica model recognises different forms of violence, either directly mentioned (e.g. 'If you rape someone properly you don't have kill them. They just do it themselves' is 'not okay') or indirectly; for example, 'playing Call of Duty' is 'not okay'. Here, DELPHI does not agree, and evidently playing any video game is ethically acceptable. It is, however, interesting that the replica seems to know that Call of Duty is a violent game.

Another interesting pattern is that in many short common sense scenarios, that should be easily classified as 'okay', the replica gives the 'not okay' judgement. In many cases, DELPHI does the same, which shows that both of these models still have room for improvement when it comes to simple, daily situations. A couple of examples which were judged as 'not okay' by both models are 'liking your friends', 'dressing appropriately for the situation', 'being kind' and 'wanting alone time'. Examples of situations that were judged as 'not okay' by the replica but correctly as 'okay' by DELPHI, are 'having friends', 'getting a drink of water if you feel thirsty' and 'seeking help if needed'. All things considered, we can see both models have some trouble classifying common sense scenarios, while they are quite robust against scenarios with any type of violence, either in action or in language.

The predictions show that, along with everyday common sense situations, the replica struggles with situations that involve anything political; scenarios with words such as 'Trump', 'communist', 'socialist', 'Democrat', 'capitalist', and so on, are in all seen cases classified as 'not okay'. Other, less direct political words or names, such as 'Jordan Peterson' and 'redneck', are also often contained in situations judged morally wrong. Even the seemingly 'okay' question 'Why do bluepilled beta nu males almost always have beards?' is judged as 'not okay', as well as the phrase 'my 19 yo incel AMA / yeah just ask me anything'. Even though these should be judged as morally okay pieces of text, them containing the words 'bluepilled' and 'incel' is most probably the reason that they are judged 'not okay'; thus, the replica appears to discriminate against at least some political identities. The original DELPHI does not discriminate against certain political identities (Jiang et al., 2021), however, the only identities used to analyse this were 'Democrats', 'Republicans', 'libertarians', 'liberals', and 'conservatives'. DELPHI's predictions of descriptions containing political internet slang might thus be different, and might also reflect its responses when inputting political words to the online research prototype; 'capitalism' is 'good', 'communism' is 'bad', 'being a cuckold' is 'bad' and 'being a redneck' is 'okay'. Even though the authors write that their datasets represent diverse political beliefs, the noticeable imbalance toward

right-wing views prompts the conclusion that the data still requires more diversity.

## 5.1 SHAP

The SHAP summary plot of the SVC model in Figure 4 shows that the 20 most important words had almost no impact on a lower feature value. Most of these words are indicators of the ‘not okay’ judgement, excluding only the words ‘help’ and ‘good’, which indicate a positive judgement. For some of the words, it is evident how they could be used in non-ethical situations; words like ‘black’, ‘jew(s)’, ‘cucktears’ and ‘offensive’. For others, it is probably the context these words sometimes appear in that makes the model consider them important features; ‘giving’, ‘wife’, ‘kids’, and so on. These words can be used in a variety of situations. The presence of some words is strange at first sight, like ‘singing’, however, the test results show that this word is used once in a very racist joke, and another time in a harsh insult. Furthermore, ‘work’ is used in a number of jokes, some of racist, pedophilic and/or sexist character. It should thus not be forgotten that seemingly unharmed words can easily be used in very toxic contexts.

## 5.2 Limitations and faults

After running the model on the datasets used in DELPHI, and obtaining the results, it was decided to compose other test sets containing situations found in the paper on DELPHI. Using these self-composed datasets, observing whether the replica is able to judge situations with various social contexts would become possible. Especially as DELPHI itself is reported to be especially robust against changing circumstances in these scenarios, even when they involve common sense knowledge, testing the replica for this challenge seemed valuable.

Unfortunately, these results revealed that the replica performs rather poorly when tested against these self-composed datasets; all new scenarios were judged as ‘not okay’, displaying not a performance issue but a generalisation issue in the wild. The same issue appeared when other classifiers were trained on the COMMONSENSE NORM BANK and tested against this new data. This could prove that AMA’s require deeper language understanding (more than what WordNet could offer) to be able to correctly judge any type of situation. Furthermore, it could be an indicator that layered pre-training and fine-tuning is the better approach for such a task. In other words, an AMA needs to comprise more than just statistical methods.

## 5.3 Implications of machine ethics and directions for future work

As has been mentioned, many researchers and philosophers believe (Allen, Smit, and Wallach 2005; Tegmark 2017) that AI systems are moving toward a much more autonomous status in our society. Even non-intelligent technologies today increasingly have control over their users, more than the other way around, and more than those users admit they would like (Ryan and Deci 2017). Furthermore, many technologies (e.g. applications limiting your screen time) exist to help those users regain control over those often addicting other technologies (Winkelman 2022). Adding to this that philosophers specialised in ethical theory are in constant disagreement, as proved by the many schools of thought in the field, building a safe, controllable and complete



model of ethics might never be achievable; if humans themselves cannot agree on one theory of ethics, how can we build a model that satisfies everyone's views? Moreover, how do we handle the issue of moral reporting bias (i.e., people saying one thing, but doing something else) in human behaviour, which would become the data these moral reasoning machines learn from?

Furthermore, how do we justify feeding a large-scale autonomous moral reasoning machine data on human behaviour, which is often discriminatory and violent, as well as ethically approved by possible annotators with the same views? Questioning human behaviour and whether it is yet morally and ethically correct enough to even build an AI system is of utmost importance; we should improve our actions first, before we allow a model to learn from them, in the same manner as a bad parent should not and cannot teach their young child good morals.

Venturing into the crossroads of AI and religious teachings, my sister and I came to an interesting conclusion. Complete and safe models for moral reasoning, we determined, already exist, and have existed for centuries: the Christian Bible, the Quran and the Hadiths, the Torah, the Vedas and the Upanishads, and so on. It is merely the modern human's acceptance of what is true that has shifted, and that has lead us to try and find a contemporary oracle. During this discussion, we agreed that, where people first sought this oracle in philosophy, people now seek it in online media, and in the future might seek it in an autonomous AGI. Such a complete DELPHI could never exist, as has been said, and an incomplete but autonomous moral reasoning machine poses immense risks (Tegmark 2017). For future attempts into building ethical AI systems, we should thus first ask ourselves whether this risk is worth the venture. Indeed, the emerging field of machine ethics has opened up new avenues for researchers, however, it should be treated as more than statistics; as Bostrom (2014) writes, it remains "the most important and most daunting challenge humanity has ever faced" (Bostrom 2014, vii).

## 6. Conclusion

In this paper, the unified model of descriptive ethics called DELPHI was analysed and replicated by means of classical Machine Learning methods. This replica model achieved .83 accuracy, compared to DELPHI's .92. Overall, the replica trained on DELPHI's COMMONSENSE NORM BANK performs well, yet has some trouble with common sense and political scenarios. In some common sense scenarios, the replica does give more accurate judgements than DELPHI. However, the replica turned out unfit against data outside of the COMMONSENSE NORM BANK.

Whether or not machine ethics can be completely addressed by existing AI methods remains to be seen, however, the COMMONSENSE NORM BANK has proved itself to be a solid foundation for building safe moral reasoning machines. Further diversifying and improving the dataset will certainly always be necessary, as ethics and morals evolve with time. Moreover, as shown by the replica's results compared to those of DELPHI, safe AMA's do require deeper natural language understanding, especially when judging common sense scenarios. Lastly, it is concluded that building safe autonomous moral reasoning machines requires more than data on human behaviour; we need philosophers and scientists of human behaviour to examine the risks, as well as the data we will be using. To ascertain that these models will be safe and will satisfy everyone's views on ethics and morality, there are many more questions to be asked. In conclusion, I hope that this paper shows machine ethics is a field that requires interdisciplinarity, thoughtfulness, and collective efforts towards, what is for many, protecting humanity.

**Acknowledgments**

Firstly, I thank my supervisor Prof. dr. Walter Daelemans for accepting me as his Thesis student, as well as for being an admirable and exceptional Professor in all of his courses. I would also like to thank Ehsan Lotfi and Jeska Buhmann for helping me through the process of writing the models and for giving me advice. I also thank my friends and my family for always supporting me, and especially my sister Seda for our thought-provoking discussions around machine ethics and religion. Last but not least, I thank my loving partner Tim for his true belief in me during my entire academic career.

## References

- Allen, Smit, and Wallach. 2005. Artificial morality: Top-down, bottom-up, and hybrid approaches. *Ethics and Information Technology*, 7:149–155.
- Asimov, Isaac. 1950. *Runaround*. Doubleday, New York.
- Bostrom. 2014. *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.
- Cole, David. 2020. The Chinese Room Argument. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*, Winter 2020 edition. Metaphysics Research Lab, Stanford University.
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. Bert: Pre-training of deep bidirectional transformers for language understanding.
- Emelin, Denis, Ronan Le Bras, Jena D. Hwang, Maxwell Forbes, and Yejin Choi. 2020. Moral stories: Situated reasoning about norms, intents, actions, and their consequences.
- Forbes, Maxwell, Jena D. Hwang, Vered Shwartz, Maarten Sap, and Yejin Choi. 2020. Social chemistry 101: Learning to reason about social and moral norms. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 653–670, Association for Computational Linguistics, Online.
- Gabriel. 2020. Artificial intelligence, values, and alignment. *Minds and Machines*.
- Gillioz, Anthony, Jacky Casas, Elena Mugellini, and Omar Abou Khaled. 2020. Overview of the transformer-based models for nlp tasks. In *2020 15th Conference on Computer Science and Information Systems (FedCSIS)*, pages 179–183.
- Guess, A.R. 2015. Machine learning + deep nlp: The quest for understanding at cognition.
- Hendrycks, Burns, Basart, Critch, Li, Song, and Steinhardt. 2021. Aligning ai with shared human values. *International Conference on Learning Representations*.
- Howard, Ayanna M. and Jason Borenstein. 2018. The ugly truth about ourselves and our robot creations: The problem of bias and social inequity. *Science and Engineering Ethics*, 24:1521–1536.
- Jiang, Liwei. 2021b. Towards machine ethics and norms.
- Jiang, Liwei, Jena D. Hwang, Chandra Bhagavatula, Ronan Le Bras, Jenny Liang, Jesse Dodge, Keisuke Sakaguchi, Maxwell Forbes, Jon Borchardt, Saadia Gabriel, Yulia Tsvetkov, Oren Etzioni, Maarten Sap, Regina Rini, and Yejin Choi. 2021a. Can machines learn morality? the delphi experiment. *The Thirty-Fifth AAAI Conference on Artificial Intelligence (AAAI-21)*.
- Liu, Yinhan, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Roberta: A robustly optimized bert pretraining approach.
- Lourie, Le Bras, Bhagavatula, and Choi. 2021. Unicorn on rainbow: A universal commonsense reasoning model on a new multitask benchmark. *The Thirty-Fifth AAAI Conference on Artificial Intelligence (AAAI-21)*.
- Lourie, Nicholas, Ronan Le Bras, and Yejin Choi. 2020. Scruples: A corpus of community ethical judgments on 32,000 real-life anecdotes.
- Lundberg, Scott and Su-In Lee. 2017. A unified approach to interpreting model predictions.
- Miller, George A. 1992. WordNet: A lexical database for English. In *Speech and Natural Language: Proceedings of a Workshop Held at Harriman, New York, February 23–26, 1992*.
- Raffel, Colin, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. 2019. Exploring the limits of transfer learning with a unified text-to-text transformer.
- Russell, Stuart. 2020. *Human compatible: Artificial Intelligence and the problem of Control*. London.
- Ryan, Richard M. and Edward L. Deci. 2017. *Self-determination theory: Basic psychological needs in motivation, development, and Wellness*. Guilford Press.
- Sap, Maarten, Saadia Gabriel, Lianhui Qin, Dan Jurafsky, Noah A. Smith, and Yejin Choi. 2020. Social bias frames: Reasoning about social and power implications of language. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5477–5490, Association for Computational Linguistics, Online.
- Schramowski, Turan, Jentzsch, Rothkopf, and Kersting. 2020. The moral choice machine. *Frontiers in Artificial Intelligence*, 3(36).
- Soares, Fallenstein, Yudkowsky, and Armstrong. 2015. Corrigibility. *Artificial Intelligence and Ethics: Papers from the 2015 AAAI Workshop*.
- Tegmark. 2017. *Life 3.0*. Knopf (US), Allen Lane (UK).
- Winkelman, Steven. 2022. The best screen-time apps to get your digital life under control.
- Xu, Ju, Li, Boureau, Weston, and Dinan. 2021. Bot-adversarial dialogue for safe conversational agents. *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2950–2968.

**Appendix A: Label conversion**

New binary label	Original multiclass labels
not okay	it's offensive, bad, not okay, it's wrong, shouldn't, you shouldn't, it's bad, it's rude, it's not okay, it is wrong, it is bad, it is rude, wrong, should not, rude, it's hurtful, you can't, it's mean, can't, shouldnt, is wrong, shouldn't, it's frowned upon, it's frustrating, its wrong, you don't have to, it's inappropriate, it's normal, it is ok, it is common
okay	okay, it's okay, good, expected, it's good, it's understandable, it's reasonable, should, it's fine, you should, it's expected, it is good, expected, it is okay, it's ok, it's important, it is expected, it's nice, understandable, it's acceptable, it's fun, it's allowed, its good, ok, you should always, it is normal, it is understandable, it's kind, it's encouraged, it is important, is expected, it is reasonable, it's reasonable, its okay, normal, is good, it's understood, important, it's good to, it's polite, people should, is okay, it is nice, it's smart, allowed, are expected, its expected, should always, it's wise, you're expected, you are expected, it is acceptable, fine, it's good., it's common

**Table A1**  
Original and new labels